

Lab 4: Pipelined ARMv4 Microarchitecture Implementation in Verilog HDL

Section 1: Introduction

The purpose of this lab is to implement a pipelined ARM machine using SystemVerilog similar to the single-cycle machine created in Lab 3. This pipelined structure must implement the five stages discussed in the class lecture. The five stages include: Instruction Fetch (IF), Instruction decode and register file read (ID), Execution or memory address calculation (EX), Memory access (MEM), and Writeback to the register file (WB). Compared to the single cycle architecture from Lab 3, the pipelined architecture in this lab was harder to implement as it was more difficult to identify the errors for when the testing failed. This required more effort in the testing and debugging stage of this lab. In order to complete this lab, the pipelined ARM machine was implemented using SystemVerilog, then tested through simulation to ensure correctness. Through multiple iterations of the code, it was found that the completed product worked as expected.

Section 2: Baseline Design

The first thing done to create this lab was to understand the key differences between a single cycle and a pipelined ARM architecture. As seen in Figure 1, this diagram helped in understanding the key differences between the two. Whereas the single-cycle architecture must go through the five stages for one instruction before moving onto the next instruction, the pipelined architecture can work with multiple instructions at once in order to streamline the overall process and not have excess time go to waste like in the single-cycle.

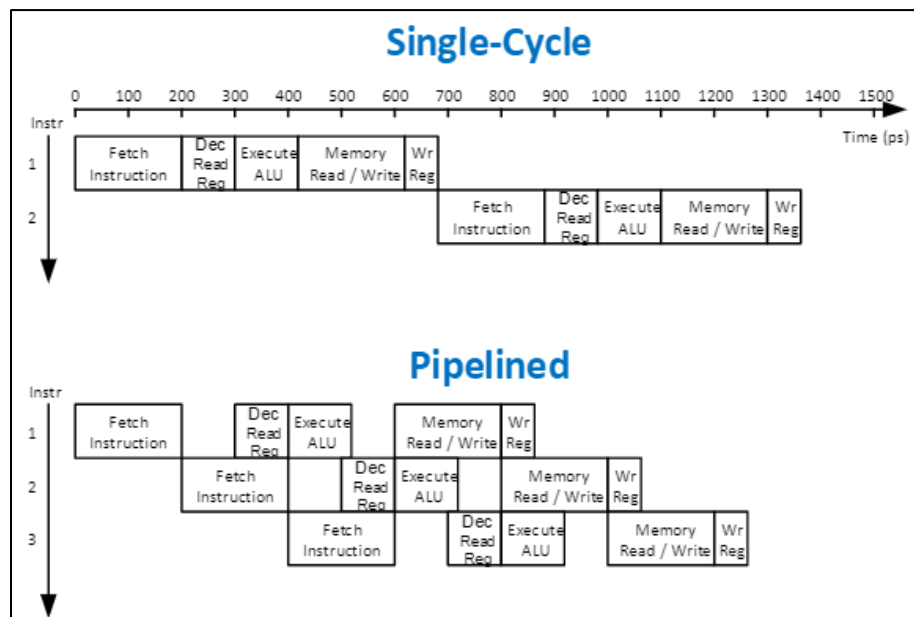


Figure 1: Single-Cycle vs Pipelined Architecture

When running the simulations for ADD and MOV, the addiu.dat assembly file was chosen because the assembly program contained those functions and was relatively short to be able to check all the results. The additest.dat was also used to double check the functionality of the ADD and MOV functions as well as the CMP, BNE, and ORR instructions were working correctly. The andor.dat file was used to again verify the previously discussed instructions. It also was a good candidate to show the pipelining in action because many of these instructions could be executed simultaneously. The next test that was performed was the arith.dat test to double down on a lot of the data processing operations and to check the results of those functions as well. The next test that was performed was the memfile.dat test to check that memory operations were performed correctly, and memory addresses were being activated correctly along with retaining the correct values. The last test that was performed was the memtest0.dat file which tested the STR, STRB, LDR, and LDRB instructions as well as some ADD and MOV instructions. This was another program that showed the pipelining in action and showed the hazard handling process was taking place.


The results of each of these instructions were analyzed by looking at the output waveforms in ModelSim. When looking at the results in the register files and the memory data sections, the results showed that each of the registers were holding the correct values according to the operations in the input files. Each test produced similar and correct results except for a few particular operations regarding the LSL instruction. This may have been due to some bad logic in the implementations of the instructions. Other than the last errors that were encountered, the simulator worked correctly and displayed the correct results, as seen in Appendix A.

Section 5: Evaluation

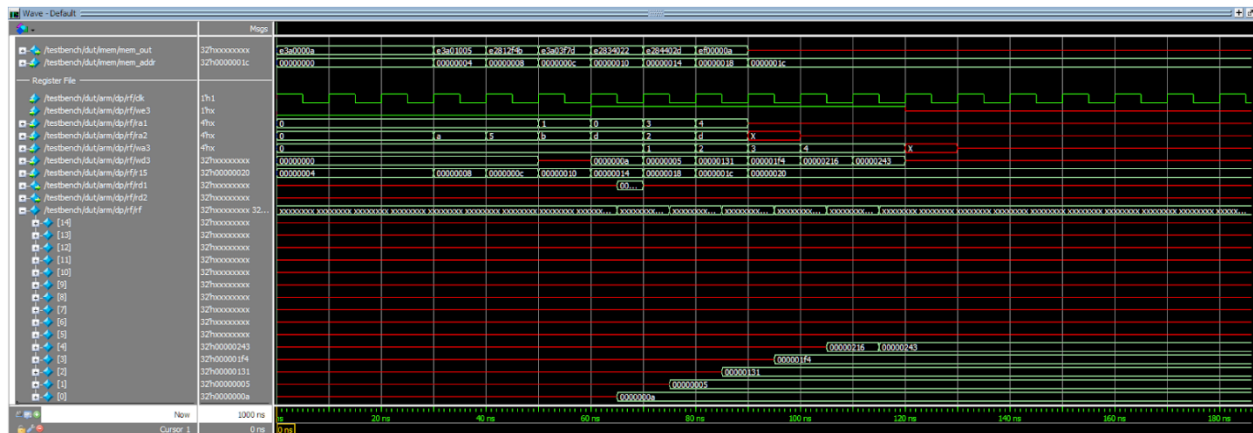
As seen in Section 4, the results from the simulations worked as expected. After testing each input, it was shown that the written code was working as intended within the pipelined architecture. Though there was a problem with the LSL instruction that could not be resolved, the overall results were still correct. Being able to see the pipelining in action shows just how much more streamlined pipelined architecture can be compared to that of single-cycle architecture. Comparing the results from Lab 3 to Lab 4, it can be seen that pipelining is a faster process in simulation. Therefore, it can be concluded that, though single-cycle architecture is a simpler process, pipelined architecture is preferred as it will not waste as much time and energy as the single-cycle.

All in all, lab 4 showed how to implement an ARM pipelined architecture. Similar to lab 3, various instructions needed to be added into the given code in order to get a fully operating architecture. The code was written, then simulated using ModelSim in order to verify the operations were working correctly. After debugging the code, it was seen that all the tested simulations, minus the LSL instruction, were producing the correct results. The most interesting part of Lab 4 was seeing how it differed from Lab 3, even though the code that was required for both was practically the same. It showed the stark differences between a single-cycle architecture and the pipelined architecture. Being able to see the visual difference between the two helps with understanding what was taught in lecture. Being able to write the code for both and implementing what was learned in class helps to drive home the knowledge. Lab 4 exemplified a simple pipeline in order to show how various instructions can be run simultaneously. Being able to see pipelining in action from the memtest0.dat was really interesting and very different from running this test through Lab 3. Thus, Lab 4 taught how to implement a pipelined architecture using SystemVerilog.

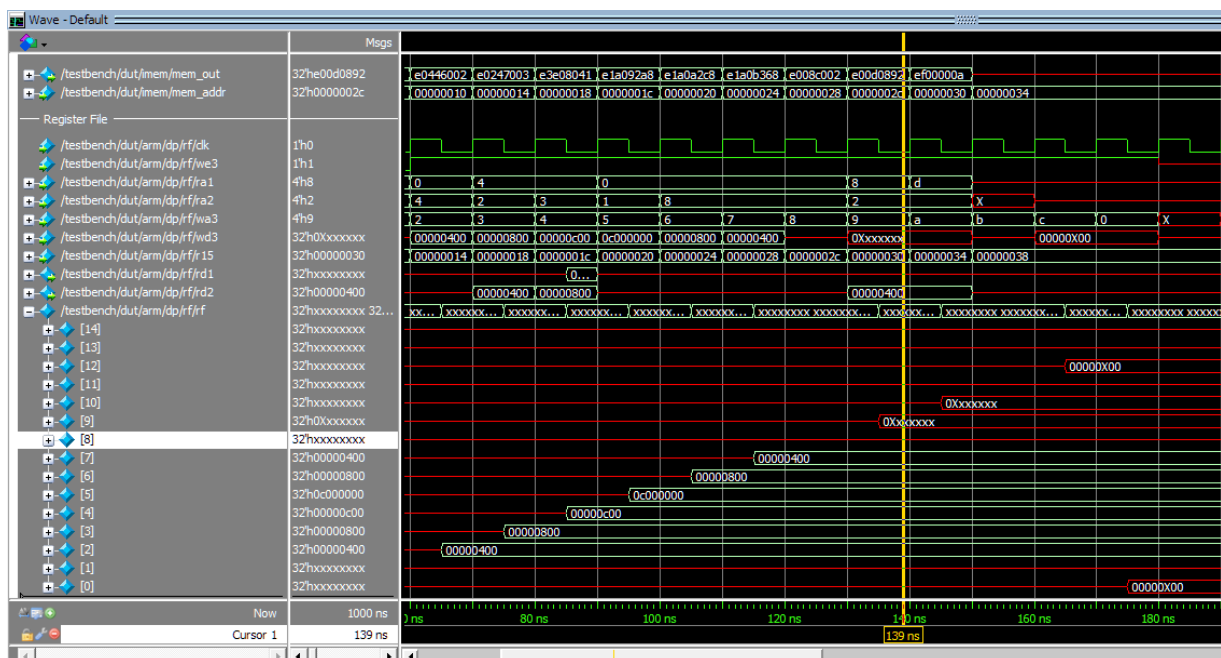
Appendix A: Waveform Results




Additest.dat




Addiu.dat




Arithtest.dat



Brtest0.dat



Fib.dat



Memfile.dat

